

# Juan Garcia

San Francisco, CA • john.doe@email.com • rendercv.com • rendercv • rendercv

---

## Welcome to RenderCV

---

RenderCV reads a CV written in a YAML file, and generates a PDF with professional typography.

Each section title is arbitrary.

You can choose any of the 9 entry types for each section.

Markdown syntax is supported everywhere. This is **bold**, *italic*, and link.

---

## Education

---

- PhD**   **Princeton University**, PhD in Computer Science – Princeton, NJ   Sept 2018 – May 2023
- Thesis: Efficient Neural Architecture Search for Resource-Constrained Deployment
  - Advisor: Prof. Sanjeev Arora
  - NSF Graduate Research Fellowship, Siebel Scholar (Class of 2022)
- BS**   **Boğaziçi University**, BS in Computer Engineering – Istanbul, Türkiye   Sept 2014 – June 2018
- GPA: 3.97/4.00, Valedictorian
  - Fulbright Scholarship recipient for Graduate Studies

---

## Experience

---

- Nexus AI**, Co-Founder & CTO – San Francisco, CA   June 2023 – present
- Built foundation model infrastructure serving 2M+ monthly API requests with 99.97% uptime
  - Raised \$18M Series A led by Sequoia Capital, with participation from a16z and Founders Fund
  - Scaled engineering team from 3 to 28 across ML research, platform, and applied AI divisions
  - Developed proprietary inference optimization reducing latency by 73% compared to baseline
- NVIDIA Research**, Research Intern – Santa Clara, CA   May 2022 – Aug 2022
- Designed sparse attention mechanism reducing transformer memory footprint by 4.2x
  - Co-authored paper accepted at NeurIPS 2022 (spotlight presentation, top 5% of submissions)
- Google DeepMind**, Research Intern – London, UK   May 2021 – Aug 2021
- Developed reinforcement learning algorithms for multi-agent coordination
  - Published research at top-tier venues with significant academic impact
  - ICML 2022 main conference paper, cited 340+ times within two years
  - NeurIPS 2022 workshop paper on emergent communication protocols
  - Invited journal extension in JMLR (2023)
- Apple ML Research**, Research Intern – Cupertino, CA   May 2020 – Aug 2020
- Created on-device neural network compression pipeline deployed across 50M+ devices
  - Filed 2 patents on efficient model quantization techniques for edge inference
- Microsoft Research**, Research Intern – Redmond, WA   May 2019 – Aug 2019
- Implemented novel self-supervised learning framework for low-resource language modeling
  - Research integrated into Azure Cognitive Services, reducing training data requirements by 60%

---

## Projects

---

- FlashInfer**   Jan 2023 – present
- Open-source library for high-performance LLM inference kernels
- Achieved 2.8x speedup over baseline attention implementations on A100 GPUs
  - Adopted by 3 major AI labs, 8,500+ GitHub stars, 200+ contributors
- NeuralPrune**   Jan 2021
- Automated neural network pruning toolkit with differentiable masks
- Reduced model size by 90% with less than 1% accuracy degradation on ImageNet
  - Featured in PyTorch ecosystem tools, 4,200+ GitHub stars

---

## Publications

---

|                                                                                                                                                                                       |           |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------|
| <b>Sparse Mixture-of-Experts at Scale: Efficient Routing for Trillion-Parameter Models</b><br><i>John Doe, Sarah Williams, David Park</i><br>10.1234/neurips.2023.1234 (NeurIPS 2023) | July 2023 |
| <b>Neural Architecture Search via Differentiable Pruning</b><br><i>James Liu, John Doe</i><br>10.1234/neurips.2022.5678 (NeurIPS 2022, Spotlight)                                     | Dec 2022  |
| <b>Multi-Agent Reinforcement Learning with Emergent Communication</b><br><i>Maria Garcia, John Doe, Tom Anderson</i><br>10.1234/icml.2022.9012 (ICML 2022)                            | July 2022 |
| <b>On-Device Model Compression via Learned Quantization</b><br><i>John Doe, Kevin Wu</i><br>10.1234/iclr.2021.3456 (ICLR 2021, Best Paper Award)                                      | May 2021  |

---

## Selected Honors

---

- MIT Technology Review 35 Under 35 Innovators (2024)
- Forbes 30 Under 30 in Enterprise Technology (2024)
- ACM Doctoral Dissertation Award Honorable Mention (2023)
- Google PhD Fellowship in Machine Learning (2020 – 2023)
- Fulbright Scholarship for Graduate Studies (2018)

---

## Skills

---

**Languages:** Python, C++, CUDA, Rust, Julia

**ML Frameworks:** PyTorch, JAX, TensorFlow, Triton, ONNX

**Infrastructure:** Kubernetes, Ray, distributed training, AWS, GCP

**Research Areas:** Neural architecture search, model compression, efficient inference, multi-agent RL

---

## Patents

---

1. Adaptive Quantization for Neural Network Inference on Edge Devices (US Patent 11,234,567)
2. Dynamic Sparsity Patterns for Efficient Transformer Attention (US Patent 11,345,678)
3. Hardware-Aware Neural Architecture Search Method (US Patent 11,456,789)

---

## Invited Talks

---

4. Scaling Laws for Efficient Inference — Stanford HAI Symposium (2024)
3. Building AI Infrastructure for the Next Decade — TechCrunch Disrupt (2024)
2. From Research to Production: Lessons in ML Systems — NeurIPS Workshop (2023)
1. Efficient Deep Learning: A Practitioner's Perspective — Google Tech Talk (2022)